

cnn

April 23, 2019

1 Home 3: Build a CNN for image recognition.

1.0.1 Name: Yao Xiao

1.1 1. Data preparation

1.1.1 1.1. Load data

```
In [2]: from keras.datasets import cifar10
import numpy as np

(x_train, y_train), (x_test, y_test) = cifar10.load_data()

print('shape of x_train: ' + str(x_train.shape))
print('shape of y_train: ' + str(y_train.shape))
print('shape of x_test: ' + str(x_test.shape))
print('shape of y_test: ' + str(y_test.shape))
print('number of classes: ' + str(np.max(y_train) - np.min(y_train) + 1))
```

```
shape of x_train: (50000, 32, 32, 3)
shape of y_train: (50000, 1)
shape of x_test: (10000, 32, 32, 3)
shape of y_test: (10000, 1)
number of classes: 10
```

1.1.2 1.2. One-hot encode the labels

In the input, a label is a scalar in $\{0, 1, \dots, 9\}$. One-hot encode transform such a scalar to a 10-dim vector. E.g., a scalar $y_{\text{train}}[j]=3$ is transformed to the vector $y_{\text{train_vec}}[j]=[0, 0, 0, 1, 0, 0, 0, 0, 0, 0]$.

1. Define a function `to_one_hot` that transforms an $n \times 1$ array to a $n \times 10$ matrix.
2. Apply the function to y_{train} and y_{test} .

```
In [3]: def to_one_hot(y, num_class=10):
        result = np.zeros((len(y), num_class))
        for i, label in enumerate(y):
```

```

        result[i,label]=1
    return result
y_train_vec = to_one_hot(y_train)
y_test_vec = to_one_hot(y_test)

print('Shape of y_train_vec: ' + str(y_train_vec.shape))
print('Shape of y_test_vec: ' + str(y_test_vec.shape))

print(y_train[0])
print(y_train_vec[0])

```

```

Shape of y_train_vec: (50000, 10)
Shape of y_test_vec: (10000, 10)
[6]
[0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]

```

Remark: the outputs should be

- Shape of y_train_vec: (50000, 10)
- Shape of y_test_vec: (10000, 10)
- [6]
- [0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]

1.1.3 1.3. Randomly partition the training set to training and validation sets

Randomly partition the 50K training samples to 2 sets: * a training set containing 40K samples * a validation set containing 10K samples

```

In [4]: rand_indices = np.random.permutation(50000)
        train_indices = rand_indices[0:40000]
        valid_indices = rand_indices[40000:50000]

        x_val = x_train[valid_indices, :]
        y_val = y_train_vec[valid_indices, :]

        x_tr = x_train[train_indices, :]
        y_tr = y_train_vec[train_indices, :]

        print('Shape of x_tr: ' + str(x_tr.shape))
        print('Shape of y_tr: ' + str(y_tr.shape))
        print('Shape of x_val: ' + str(x_val.shape))
        print('Shape of y_val: ' + str(y_val.shape))

```

```

Shape of x_tr: (40000, 32, 32, 3)
Shape of y_tr: (40000, 10)
Shape of x_val: (10000, 32, 32, 3)
Shape of y_val: (10000, 10)

```

1.2 2. Build a CNN and tune its hyper-parameters

1. Build a convolutional neural network model
2. Use the validation data to tune the hyper-parameters (e.g., network structure, and optimization algorithm)
3. Try to achieve a validation accuracy as high as possible.

```
In [4]: from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
        from keras.models import Sequential
```

```
model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(32, 32, 3)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(10, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_1 (Dense)	(None, 128)	524416
dense_2 (Dense)	(None, 10)	1290

Total params: 545,098
Trainable params: 545,098
Non-trainable params: 0

```
In [5]: from keras import optimizers
```

```
learning_rate = 1E-5 # to be tuned!

model.compile(loss='categorical_crossentropy',
```

```
optimizer=optimizers.RMSprop(lr=learning_rate),
metrics=['acc'])
```

```
In [6]: history = model.fit(x_tr, y_tr, batch_size=32, epochs=10, validation_data=(x_val, y_val))
```

Train on 40000 samples, validate on 10000 samples


```
Epoch 1/10
40000/40000 [=====] - 36s 905us/step - loss: 8.2007 - acc: 0.2073 - val_loss: 8.2007 - val_acc: 0.2073
Epoch 2/10
40000/40000 [=====] - 39s 964us/step - loss: 4.9410 - acc: 0.2873 - val_loss: 4.9410 - val_acc: 0.2873
Epoch 3/10
40000/40000 [=====] - 39s 970us/step - loss: 3.2000 - acc: 0.3265 - val_loss: 3.2000 - val_acc: 0.3265
Epoch 4/10
40000/40000 [=====] - 39s 972us/step - loss: 2.5738 - acc: 0.3583 - val_loss: 2.5738 - val_acc: 0.3583
Epoch 5/10
40000/40000 [=====] - 39s 975us/step - loss: 2.2260 - acc: 0.3864 - val_loss: 2.2260 - val_acc: 0.3864
Epoch 6/10
40000/40000 [=====] - 39s 982us/step - loss: 1.9925 - acc: 0.4120 - val_loss: 1.9925 - val_acc: 0.4120
Epoch 7/10
40000/40000 [=====] - 39s 986us/step - loss: 1.8287 - acc: 0.4350 - val_loss: 1.8287 - val_acc: 0.4350
Epoch 8/10
40000/40000 [=====] - 40s 988us/step - loss: 1.7027 - acc: 0.4558 - val_loss: 1.7027 - val_acc: 0.4558
Epoch 9/10
40000/40000 [=====] - 40s 990us/step - loss: 1.6063 - acc: 0.4773 - val_loss: 1.6063 - val_acc: 0.4773
Epoch 10/10
40000/40000 [=====] - 40s 993us/step - loss: 1.5326 - acc: 0.4964 - val_loss: 1.5326 - val_acc: 0.4964
```

```
In [7]: import matplotlib.pyplot as plt
        %matplotlib inline
```

```
acc = history.history['acc']
val_acc = history.history['val_acc']

epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'r', label='Validation acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



1.3 3. Train (again) and evaluate the model

- To this end, we have found the "best" hyper-parameters.
- Now, fix the hyper-parameters and train the network on the entire training set (all the 50K training samples)
- Evaluate the model on the test set.

1.3.1 3.1. Train the model on the entire training set

Why? Previously, we used 40K samples for training; we wasted 10K samples for the sake of hyper-parameter tuning. Now we already know the hyper-parameters, so why not using all the 50K samples for training?

```
In [28]: from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Activation
         from keras.models import Sequential
         from keras import layers
         model1 = Sequential()
         model1.add(layers.Conv2D(32, (3, 3), padding='same', input_shape=(32, 32, 3)))
         model1.add(layers.BatchNormalization())
         model1.add(Activation('relu'))
         model1.add(layers.MaxPooling2D((2, 2)))
         model1.add(layers.Conv2D(64, (3, 3), padding='same'))
         model1.add(layers.BatchNormalization())
         model1.add(Activation('relu'))
         model1.add(layers.MaxPooling2D((2, 2)))
```

```

model1.add(layers.Conv2D(128, (3, 3), padding='same'))
model1.add(layers.BatchNormalization())
model1.add(Activation('relu'))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Flatten())
model1.add(layers.Dropout(0.5))
model1.add(layers.Dense(512))
model1.add(layers.BatchNormalization())
model1.add(Activation('relu'))
model1.add(layers.Dense(10, activation='softmax'))

model1.summary()

```

Layer (type)	Output Shape	Param #
conv2d_32 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_29 (Batch Normalization)	(None, 32, 32, 32)	128
activation_29 (Activation)	(None, 32, 32, 32)	0
max_pooling2d_30 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_33 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_30 (Batch Normalization)	(None, 16, 16, 64)	256
activation_30 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_31 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_34 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_31 (Batch Normalization)	(None, 8, 8, 128)	512
activation_31 (Activation)	(None, 8, 8, 128)	0
max_pooling2d_32 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_10 (Flatten)	(None, 2048)	0
dropout_9 (Dropout)	(None, 2048)	0
dense_16 (Dense)	(None, 512)	1049088
batch_normalization_32 (Batch Normalization)	(None, 512)	2048

```

activation_32 (Activation)      (None, 512)                  0
-----
dense_17 (Dense)                (None, 10)                   5130
=====
Total params: 1,150,410
Trainable params: 1,148,938
Non-trainable params: 1,472
-----

```

```
In [29]: from keras import optimizers
```

```

learning_rate = 3E-4 # to be tuned!
model1.compile(loss='categorical_crossentropy', optimizer=optimizers.RMSprop(lr=learn

```

```
In [30]: history = model1.fit(x_train, y_train_vec, batch_size=128, epochs=50)
```

```

Epoch 1/50
50000/50000 [=====] - 132s 3ms/step - loss: 1.4575 - acc: 0.4771
Epoch 2/50
50000/50000 [=====] - 129s 3ms/step - loss: 1.1073 - acc: 0.6070
Epoch 3/50
50000/50000 [=====] - 129s 3ms/step - loss: 0.9777 - acc: 0.6540
Epoch 4/50
50000/50000 [=====] - 127s 3ms/step - loss: 0.8880 - acc: 0.6880
Epoch 5/50
50000/50000 [=====] - 118s 2ms/step - loss: 0.8225 - acc: 0.7113
Epoch 6/50
50000/50000 [=====] - 119s 2ms/step - loss: 0.7734 - acc: 0.7268
Epoch 7/50
50000/50000 [=====] - 129s 3ms/step - loss: 0.7264 - acc: 0.7449
Epoch 8/50
50000/50000 [=====] - 131s 3ms/step - loss: 0.6924 - acc: 0.7574
Epoch 9/50
50000/50000 [=====] - 131s 3ms/step - loss: 0.6564 - acc: 0.7680
Epoch 10/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.6276 - acc: 0.7795
Epoch 11/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.5939 - acc: 0.7916
Epoch 12/50
50000/50000 [=====] - 131s 3ms/step - loss: 0.5707 - acc: 0.7979
Epoch 13/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.5450 - acc: 0.8064
Epoch 14/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.5189 - acc: 0.8189
Epoch 15/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.5037 - acc: 0.8240
Epoch 16/50

```

50000/50000 [=====] - 131s 3ms/step - loss: 0.4829 - acc: 0.8293
 Epoch 17/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.4656 - acc: 0.8364
 Epoch 18/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.4479 - acc: 0.8422
 Epoch 19/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.4309 - acc: 0.8477
 Epoch 20/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.4112 - acc: 0.8557
 Epoch 21/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3950 - acc: 0.8595
 Epoch 22/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3857 - acc: 0.8655
 Epoch 23/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3663 - acc: 0.8716
 Epoch 24/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3550 - acc: 0.8746
 Epoch 25/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3447 - acc: 0.8797
 Epoch 26/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3318 - acc: 0.8845
 Epoch 27/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3163 - acc: 0.8890
 Epoch 28/50
 50000/50000 [=====] - 131s 3ms/step - loss: 0.3040 - acc: 0.8943
 Epoch 29/50
 50000/50000 [=====] - 132s 3ms/step - loss: 0.2946 - acc: 0.8971
 Epoch 30/50
 50000/50000 [=====] - 132s 3ms/step - loss: 0.2850 - acc: 0.9015
 Epoch 31/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2712 - acc: 0.9059
 Epoch 32/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2618 - acc: 0.9080
 Epoch 33/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2536 - acc: 0.9108
 Epoch 34/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2502 - acc: 0.9133
 Epoch 35/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2386 - acc: 0.9182
 Epoch 36/50
 50000/50000 [=====] - 119s 2ms/step - loss: 0.2310 - acc: 0.9193
 Epoch 37/50
 50000/50000 [=====] - 123s 2ms/step - loss: 0.2251 - acc: 0.9211
 Epoch 38/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2164 - acc: 0.9245
 Epoch 39/50
 50000/50000 [=====] - 130s 3ms/step - loss: 0.2133 - acc: 0.9263
 Epoch 40/50


```

50000/50000 [=====] - 130s 3ms/step - loss: 0.2024 - acc: 0.9292
Epoch 41/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.1999 - acc: 0.9303
Epoch 42/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.1955 - acc: 0.9320
Epoch 43/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.1895 - acc: 0.9334
Epoch 44/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.1812 - acc: 0.9374
Epoch 45/50
50000/50000 [=====] - 132s 3ms/step - loss: 0.1753 - acc: 0.9394
Epoch 46/50
50000/50000 [=====] - 132s 3ms/step - loss: 0.1729 - acc: 0.9402
Epoch 47/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.1651 - acc: 0.9431
Epoch 48/50
50000/50000 [=====] - 130s 3ms/step - loss: 0.1638 - acc: 0.9433
Epoch 49/50
50000/50000 [=====] - 190s 4ms/step - loss: 0.1593 - acc: 0.9452
Epoch 50/50
50000/50000 [=====] - 132s 3ms/step - loss: 0.1594 - acc: 0.9447

```

1.3.2 3.2. Evaluate the model on the test set

```

In [32]: loss_and_acc = model1.evaluate(x_test, y_test_vec)
         print('loss = ' + str(loss_and_acc[0]))
         print('accuracy = ' + str(loss_and_acc[1]))

10000/10000 [=====] - 5s 471us/step
loss = 0.6779227051258088
accuracy = 0.8011

```

1.4 Data augmentation

```

In [9]: from keras.preprocessing.image import ImageDataGenerator

        datagen_train = ImageDataGenerator(
            width_shift_range = 0.1,
            height_shift_range = 0.1,
            horizontal_flip = True)

        datagen_train.fit(x_train)

In [21]: from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Activation
         from keras.models import Sequential
         from keras import layers
         model1 = Sequential()

```

```

model1.add(layers.Conv2D(32, (3, 3), padding='same', input_shape=(32, 32, 3)))
model1.add(layers.BatchNormalization())
model1.add(Activation('relu'))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Conv2D(64, (3, 3), padding='same'))
model1.add(layers.BatchNormalization())
model1.add(Activation('relu'))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Conv2D(128, (3, 3), padding='same'))
model1.add(layers.BatchNormalization())
model1.add(Activation('relu'))
model1.add(layers.MaxPooling2D((2, 2)))
model1.add(layers.Flatten())
model1.add(layers.Dropout(0.5))
model1.add(layers.Dense(512))
model1.add(layers.BatchNormalization())
model1.add(Activation('relu'))
model1.add(layers.Dense(10, activation='softmax'))

model1.summary()

```

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_9 (Batch Normalization)	(None, 32, 32, 32)	128
activation_9 (Activation)	(None, 32, 32, 32)	0
max_pooling2d_13 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_14 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_10 (Batch Normalization)	(None, 16, 16, 64)	256
activation_10 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_14 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_15 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_11 (Batch Normalization)	(None, 8, 8, 128)	512
activation_11 (Activation)	(None, 8, 8, 128)	0
max_pooling2d_15 (MaxPooling2D)	(None, 4, 4, 128)	0

flatten_5 (Flatten)	(None, 2048)	0

dropout_7 (Dropout)	(None, 2048)	0

dense_8 (Dense)	(None, 512)	1049088

batch_normalization_12 (Batch Normalization)	(None, 512)	2048

activation_12 (Activation)	(None, 512)	0

dense_9 (Dense)	(None, 10)	5130
=====		
Total params: 1,150,410		
Trainable params: 1,148,938		
Non-trainable params: 1,472		

```
In [22]: model1.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

```
In [17]: from keras.callbacks import ModelCheckpoint
```

```
batch_size = 32
checkpoint = ModelCheckpoint(filepath='MLP.weights.best.hdf5', verbose=1, save_best_only=True)
#history = model2.fit(x_train, y_train_vec, batch_size=128, epochs=10)
```

```
In [28]: model1.fit_generator(datagen_train.flow(x_train, y_train_vec, batch_size=32),
                             steps_per_epoch=x_train.shape[0] // batch_size,
                             epochs = 25,
                             verbose=2,
                             callbacks=[checkpoint],
                             validation_data=(x_test, y_test_vec),
                             validation_steps=x_test.shape[0] // batch_size)
```

Epoch 1/25

```
- 156s - loss: 0.7768 - acc: 0.7360 - val_loss: 0.7090 - val_acc: 0.7625
```

Epoch 00001: val_loss improved from 0.71476 to 0.70901, saving model to MLP.weights.best.hdf5

Epoch 2/25

```
- 157s - loss: 0.7601 - acc: 0.7439 - val_loss: 0.7014 - val_acc: 0.7745
```

Epoch 00002: val_loss improved from 0.70901 to 0.70140, saving model to MLP.weights.best.hdf5

Epoch 3/25

```
- 155s - loss: 0.7428 - acc: 0.7480 - val_loss: 1.3019 - val_acc: 0.6793
```

Epoch 00003: val_loss did not improve from 0.70140

Epoch 4/25

```
- 153s - loss: 0.7327 - acc: 0.7529 - val_loss: 0.7865 - val_acc: 0.7423
```

Epoch 00004: val_loss did not improve from 0.70140

Epoch 5/25

- 153s - loss: 0.7190 - acc: 0.7571 - val_loss: 0.8275 - val_acc: 0.7400

Epoch 00005: val_loss did not improve from 0.70140

Epoch 6/25

- 154s - loss: 0.7077 - acc: 0.7611 - val_loss: 0.8883 - val_acc: 0.7056

Epoch 00006: val_loss did not improve from 0.70140

Epoch 7/25

- 153s - loss: 0.7028 - acc: 0.7643 - val_loss: 0.6701 - val_acc: 0.7832

Epoch 00007: val_loss improved from 0.70140 to 0.67005, saving model to MLP.weights.best.hdf5

Epoch 8/25

- 148s - loss: 0.6931 - acc: 0.7665 - val_loss: 0.6966 - val_acc: 0.7820

Epoch 00008: val_loss did not improve from 0.67005

Epoch 9/25

- 154s - loss: 0.6881 - acc: 0.7696 - val_loss: 0.8197 - val_acc: 0.7315

Epoch 00009: val_loss did not improve from 0.67005

Epoch 10/25

- 157s - loss: 0.6801 - acc: 0.7720 - val_loss: 0.7311 - val_acc: 0.7605

Epoch 00010: val_loss did not improve from 0.67005

Epoch 11/25

- 154s - loss: 0.6689 - acc: 0.7749 - val_loss: 0.6025 - val_acc: 0.8047

Epoch 00011: val_loss improved from 0.67005 to 0.60250, saving model to MLP.weights.best.hdf5

Epoch 12/25

- 162s - loss: 0.6627 - acc: 0.7766 - val_loss: 0.6843 - val_acc: 0.7736

Epoch 00012: val_loss did not improve from 0.60250

Epoch 13/25

- 159s - loss: 0.6610 - acc: 0.7790 - val_loss: 0.8087 - val_acc: 0.7282

Epoch 00013: val_loss did not improve from 0.60250

Epoch 14/25

- 149s - loss: 0.6517 - acc: 0.7812 - val_loss: 0.6933 - val_acc: 0.7773

Epoch 00014: val_loss did not improve from 0.60250

Epoch 15/25

- 148s - loss: 0.6467 - acc: 0.7826 - val_loss: 0.5933 - val_acc: 0.8062

Epoch 00015: val_loss improved from 0.60250 to 0.59333, saving model to MLP.weights.best.hdf5

Epoch 16/25

- 152s - loss: 0.6402 - acc: 0.7859 - val_loss: 0.6885 - val_acc: 0.7700

```

Epoch 00016: val_loss did not improve from 0.59333
Epoch 17/25
- 153s - loss: 0.6360 - acc: 0.7876 - val_loss: 0.7582 - val_acc: 0.7878

Epoch 00017: val_loss did not improve from 0.59333
Epoch 18/25
- 151s - loss: 0.6338 - acc: 0.7885 - val_loss: 0.6813 - val_acc: 0.7894

Epoch 00018: val_loss did not improve from 0.59333
Epoch 19/25
- 152s - loss: 0.6293 - acc: 0.7872 - val_loss: 0.5339 - val_acc: 0.8263

Epoch 00019: val_loss improved from 0.59333 to 0.53386, saving model to MLP.weights.best.hdf5
Epoch 20/25
- 149s - loss: 0.6225 - acc: 0.7912 - val_loss: 0.7426 - val_acc: 0.7630

Epoch 00020: val_loss did not improve from 0.53386
Epoch 21/25
- 154s - loss: 0.6190 - acc: 0.7937 - val_loss: 0.6205 - val_acc: 0.8052

Epoch 00021: val_loss did not improve from 0.53386
Epoch 22/25
- 152s - loss: 0.6120 - acc: 0.7960 - val_loss: 0.6579 - val_acc: 0.8069

Epoch 00022: val_loss did not improve from 0.53386
Epoch 23/25
- 152s - loss: 0.6109 - acc: 0.7956 - val_loss: 0.6266 - val_acc: 0.8027

Epoch 00023: val_loss did not improve from 0.53386
Epoch 24/25
- 155s - loss: 0.6071 - acc: 0.7969 - val_loss: 0.5392 - val_acc: 0.8273

Epoch 00024: val_loss did not improve from 0.53386
Epoch 25/25
- 162s - loss: 0.6021 - acc: 0.7992 - val_loss: 0.6713 - val_acc: 0.7905

Epoch 00025: val_loss did not improve from 0.53386

```

```
Out[28]: <keras.callbacks.History at 0xb3c793f98>
```

```

In [30]: model1.load_weights('MLP.weights.best.hdf5')
         loss_and_acc = model1.evaluate(x_test, y_test_vec)
         print('loss = ' + str(loss_and_acc[0]))
         print('accuracy = ' + str(loss_and_acc[1]))

```

```

10000/10000 [=====] - 5s 499us/step
loss = 0.5338634187221527

```

```
accuracy = 0.8263
```